

## 1. Applications for Different RNN Architectures

### Sequence-to-Sequence RNN:

Takes a sequence as input and produces a sequence as output.

Examples:

- Machine Translation (e.g., English to French)
- Text Summarization
- Dialogue Systems
- Speech Recognition

### Sequence-to-Vector RNN:

Takes a sequence input and produces a fixed-size vector.

Examples:

- Sentiment Analysis
- Document Classification
- Video encoding before classification

### Vector-to-Sequence RNN:

Takes a fixed-size vector input and produces a sequence.

Examples:

- Image Captioning
- Text Generation from features (e.g., chatbot replies)
- Speech synthesis from text

---

## 2. Why Use Encoder–Decoder RNNs Instead of Plain Sequence-to-Sequence RNNs for Translation?

Encoder–decoder RNNs separate the model into two parts:

- Encoder compresses the input sequence into a fixed-size context vector.
- Decoder generates the output sequence from this context.

#### Advantages:

- More flexible than a single sequence-to-sequence model.
  - Handles variable-length inputs and outputs.
  - Easily extendable with attention mechanisms for better context handling.
- 

### 3. Combining CNN and RNN for Video Classification

- Use a CNN to extract features from each video frame.
- Feed the sequence of feature vectors into an RNN to capture temporal information.
- Output could be a class label or sequence of actions.

#### Example Flow:

Video → CNN (frame features) → RNN → Class Label (e.g., "walking", "running")

---

### 4. Advantages of Using `dynamic_rnn()` Over `static_rnn()`

- Memory Efficient: Dynamic RNN builds the computation graph at runtime, saving memory.
  - Faster: Especially for variable-length sequences.
  - Easier to Handle Variable Sequence Lengths: Automatically stops unrolling when sequence ends.
- 

### 5. Dealing with Variable-Length Sequences

#### Input Sequences:

- Padding: Pad shorter sequences to the max length in the batch.
- Masking: Use masks or specify sequence lengths so the model ignores padding during training.

### Output Sequences:

- Use decoder with attention or teacher forcing.
  - Stop generation using a special end-of-sequence (EOS) token.
- 

## 6. Common Way to Distribute RNN Training Across Multiple GPUs

- **Data Parallelism:** Split each batch of data across GPUs. Each GPU computes gradients on its subset. Then, gradients are averaged and weights updated.
  - Use frameworks like TensorFlow's **MirroredStrategy** or PyTorch's **DataParallel** or **DistributedDataParallel**.
  - You can also split model layers across GPUs for model parallelism, though it's less common for RNNs.
-